

# Introducing Convolutional Neural Networks

---



**Janani Ravi**

CO-FOUNDER, LOONYCORN

[www.loonycorn.com](http://www.loonycorn.com)

# Overview

**Intuition behind Convolutional Neural Networks (CNNs)**

**Convolution layers and feature maps**

**Pooling layers to subsample inputs**

**Typical CNN architecture**

# How Do We See?

---

# Viewing an Image



**All neurons in the eye don't see the  
entire image**




# Viewing an Image



**Each neuron has its own local receptive field**


# Viewing an Image



**It reacts only to visual stimuli  
located in its receptive field**




# Viewing an Image



Some neurons react to more complex patterns  
that are **combinations** of lower level patterns

# Neural Networks



Sounds like a classic neural network  
problem



# Two Kinds of Layers in CNNs

## Convolution

Local receptive field

## Pooling

Subsampling of inputs

# Convolution

---

# Two Kinds of Layers in CNNs

## Convolution

Local receptive field

## Pooling

Subsampling of inputs

# Convolution

In this context, a sliding window function applied to a matrix



# Convolution

In this context, a sliding window function applied to  
a matrix



e.g. a matrix of pixels  
representing an image

# Convolution

In this context, a sliding window **function** applied to a matrix

Often called a kernel or filter




# Convolution

In this context, a sliding window function applied to a matrix



Kernel is applied element-wise  
in sliding-window fashion


# Representing Images as Matrices



= 784 pixels



# Representing Images as Matrices




= 36 pixels

# Representing Images

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix




Kernel

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



Kernel

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



x1	x0	x1
x0	x1	x0
x1	x0	x1

4

4			
1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	1.8

Convolution  
Result



# Convolution

$0_{x1}$	$x0$	$0_{x1}$	0	0	0
$x0$	$0.8_{x1}$	$x0$	0.3	0.6	0
$0.2_{x1}$	$x0$	$0_{x1}$	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix




Convolution  
Result

# Convolution

0 <sub>x1</sub>	x0	0 <sub>x1</sub>	0	0	0
x0	0.8 <sub>x1</sub>	x0	0.3	0.6	0
0.2 <sub>x1</sub>	x0	0 <sub>x1</sub>	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1			

Convolution  
Result

# Convolution

0	0 <sub>x1</sub>	x0	0 <sub>x1</sub>	0	0
0.2	x0	0 <sub>x1</sub>	x0	0.6	0
0.2	0.9 <sub>x1</sub>	x0	0.3 <sub>x1</sub>	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1			

Convolution  
Result

# Convolution

0	0 <sub>x1</sub>	x0	0 <sub>x1</sub>	0	0
0.2	x0	0 <sub>x1</sub>	x0	0.6	0
0.2	0.9 <sub>x1</sub>	x0	0.3 <sub>x1</sub>	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2		

Convolution  
Result



# Convolution

0	0	0 <sub>x1</sub>	x0	0 <sub>x1</sub>	0
0.2	0.8	x0	0.3 <sub>x1</sub>	x0	0
0.2	0.9	0 <sub>x1</sub>	x0	0.8 <sub>x1</sub>	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2		

Convolution  
Result

# Convolution

0	0	0 <sub>x1</sub>	x0	0 <sub>x1</sub>	0
0.2	0.8	x0	0.3 <sub>x1</sub>	x0	0
0.2	0.9	0 <sub>x1</sub>	x0	0.8 <sub>x1</sub>	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	

Convolution  
Result

# Convolution

0	0	0	0 <sub>x1</sub>	x0	0 <sub>x1</sub>
0.2	0.8	0	x0	0.6 <sub>x1</sub>	x0
0.2	0.9	0	0.3 <sub>x1</sub>	x0	0 <sub>x1</sub>
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	

Convolution  
Result

# Convolution

0	0	0	0 <sub>x1</sub>	x0	0 <sub>x1</sub>
0.2	0.8	0	x0	0.6 <sub>x1</sub>	x0
0.2	0.9	0	0.3 <sub>x1</sub>	x0	0 <sub>x1</sub>
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2 <sub>x1</sub>	<b>x0</b>	0 <sub>x1</sub>	0.3	0.6	0
<b>x0</b>	0.9 <sub>x1</sub>	<b>x0</b>	0.3	0.8	0
0.3 <sub>x1</sub>	<b>x0</b>	0.7 <sub>x1</sub>	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2 <sub>x1</sub>	<b>x0</b>	0 <sub>x1</sub>	0.3	0.6	0
<b>x0</b>	0.9 <sub>x1</sub>	<b>x0</b>	0.3	0.8	0
0.3 <sub>x1</sub>	<b>x0</b>	0.7 <sub>x1</sub>	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9			

Convolution  
Result



# Convolution

0	0	0	0	0	0
0.2	0.8 <sub>x1</sub>	x0	0.3 <sub>x1</sub>	0.6	0
0.2	x0	0 <sub>x1</sub>	x0	0.8	0
0.3	0.8 <sub>x1</sub>	x0	0.8 <sub>x1</sub>	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9			

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8 <sub>x1</sub>	<b>x0</b>	0.3 <sub>x1</sub>	0.6	0
0.2	<b>x0</b>	0 <sub>x1</sub>	<b>x0</b>	0.8	0
0.3	0.8 <sub>x1</sub>	<b>x0</b>	0.8 <sub>x1</sub>	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7		

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0 <sub>x1</sub>	<b>x0</b>	0.6 <sub>x1</sub>	0
0.2	0.9	<b>x0</b>	0.3 <sub>x1</sub>	<b>x0</b>	0
0.3	0.8	0.7 <sub>x1</sub>	<b>x0</b>	0.9 <sub>x1</sub>	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7		

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0 <sub>x1</sub>	<b>x0</b>	0.6 <sub>x1</sub>	0
0.2	0.9	<b>x0</b>	0.3 <sub>x1</sub>	<b>x0</b>	0
0.3	0.8	0.7 <sub>x1</sub>	<b>x0</b>	0.9 <sub>x1</sub>	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3 <sub>x1</sub>	x0	0 <sub>x1</sub>
0.2	0.9	0	x0	0.8 <sub>x1</sub>	x0
0.3	0.8	0.7	0.8 <sub>x1</sub>	x0	0 <sub>x1</sub>
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3 <sub>x1</sub>	x0	0 <sub>x1</sub>
0.2	0.9	0	x0	0.8 <sub>x1</sub>	x0
0.3	0.8	0.7	0.8 <sub>x1</sub>	x0	0 <sub>x1</sub>
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2 <sub>x1</sub>	<b>x0</b>	0 <sub>x1</sub>	0.3	0.8	0
<b>x0</b>	0.8 <sub>x1</sub>	<b>x0</b>	0.8	0.9	0
0 <sub>x1</sub>	<b>x0</b>	0 <sub>x1</sub>	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2 <sub>x1</sub>	<b>x0</b>	0 <sub>x1</sub>	0.3	0.8	0
<b>x0</b>	0.8 <sub>x1</sub>	<b>x0</b>	0.8	0.9	0
0 <sub>x1</sub>	<b>x0</b>	0 <sub>x1</sub>	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0			

Convolution  
Result



# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	<b>0.9<sub>x1</sub></b>	<b>x0</b>	<b>0.3<sub>x1</sub></b>	0.8	0
0.3	<b>x0</b>	<b>0.7<sub>x1</sub></b>	<b>x0</b>	0.9	0
0	<b>0<sub>x1</sub></b>	<b>x0</b>	<b>0.2<sub>x1</sub></b>	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0			

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	<b>0.9<sub>x1</sub></b>	<b>x0</b>	<b>0.3<sub>x1</sub></b>	0.8	0
0.3	<b>x0</b>	<b>0.7<sub>x1</sub></b>	<b>x0</b>	0.9	0
0	<b>0<sub>x1</sub></b>	<b>x0</b>	<b>0.2<sub>x1</sub></b>	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	<b>2.1</b>		

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0 <sub>x1</sub>	<b>x0</b>	0.8 <sub>x1</sub>	0
0.3	0.8	<b>x0</b>	0.8 <sub>x1</sub>	<b>x0</b>	0
0	0	0 <sub>x1</sub>	<b>x0</b>	0.8 <sub>x1</sub>	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1		

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0 <sub>x1</sub>	<b>x0</b>	0.8 <sub>x1</sub>	0
0.3	0.8	<b>x0</b>	0.8 <sub>x1</sub>	<b>x0</b>	0
0	0	0 <sub>x1</sub>	<b>x0</b>	0.8 <sub>x1</sub>	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3 <sub>x1</sub>	x0	0 <sub>x1</sub>
0.3	0.8	0.7	x0	0.9 <sub>x1</sub>	x0
0	0	0	0.2 <sub>x1</sub>	x0	0 <sub>x1</sub>
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3 <sub>x1</sub>	x0	0 <sub>x1</sub>
0.3	0.8	0.7	x0	0.9 <sub>x1</sub>	x0
0	0	0	0.2 <sub>x1</sub>	x0	0 <sub>x1</sub>
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3 <sub>x1</sub>	<b>x0</b>	0.7 <sub>x1</sub>	0.8	0.9	0
<b>x0</b>	0 <sub>x1</sub>	<b>x0</b>	0.2	0.8	0
0 <sub>x1</sub>	<b>x0</b>	0 <sub>x1</sub>	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3 <sub>x1</sub>	<b>x0</b>	0.7 <sub>x1</sub>	0.8	0.9	0
<b>x0</b>	0 <sub>x1</sub>	<b>x0</b>	0.2	0.8	0
0 <sub>x1</sub>	<b>x0</b>	0 <sub>x1</sub>	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0			

Convolution  
Result



# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	<b>0.8<sub>x1</sub></b>	<b>x0</b>	<b>0.8<sub>x1</sub></b>	0.9	0
0	<b>x0</b>	<b>0<sub>x1</sub></b>	<b>x0</b>	0.8	0
0	<b>0<sub>x1</sub></b>	<b>x0</b>	<b>0.2<sub>x1</sub></b>	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0			

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8 <sub>x1</sub>	<b>x0</b>	0.8 <sub>x1</sub>	0.9	0
0	<b>x0</b>	0 <sub>x1</sub>	<b>x0</b>	0.8	0
0	0 <sub>x1</sub>	<b>x0</b>	0.2 <sub>x1</sub>	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8		

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	<b>0.7<sub>x1</sub></b>	<b>x0</b>	<b>0.9<sub>x1</sub></b>	0
0	0	<b>x0</b>	<b>0.2<sub>x1</sub></b>	<b>x0</b>	0
0	0	<b>0<sub>x1</sub></b>	<b>x0</b>	<b>0.2<sub>x1</sub></b>	0

Matrix



<b>1</b>	<b>1.2</b>	<b>1.1</b>	<b>0.9</b>
<b>1.9</b>	<b>2.7</b>	<b>2.5</b>	<b>1.9</b>
<b>1.0</b>	<b>2.1</b>	<b>2.4</b>	<b>1.4</b>
<b>1.0</b>	<b>1.8</b>		

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7 <sub>x1</sub>	x0	0.9 <sub>x1</sub>	0
0	0	x0	0.2 <sub>x1</sub>	x0	0
0	0	0 <sub>x1</sub>	x0	0.2 <sub>x1</sub>	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	<b>0.8<sub>x1</sub></b>	<b>x0</b>	<b>0<sub>x1</sub></b>
0	0	0	<b>x0</b>	<b>0.8<sub>x1</sub></b>	<b>x0</b>
0	0	0	<b>0.2<sub>x1</sub></b>	<b>x0</b>	<b>0<sub>x1</sub></b>

Matrix



<b>1</b>	<b>1.2</b>	<b>1.1</b>	<b>0.9</b>
<b>1.9</b>	<b>2.7</b>	<b>2.5</b>	<b>1.9</b>
<b>1.0</b>	<b>2.1</b>	<b>2.4</b>	<b>1.4</b>
<b>1.0</b>	<b>1.8</b>	<b>2.0</b>	

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8 <sub>x1</sub>	<b>x0</b>	0 <sub>x1</sub>
0	0	0	<b>x0</b>	0.8 <sub>x1</sub>	<b>x0</b>
0	0	0	0.2 <sub>x1</sub>	<b>x0</b>	0 <sub>x1</sub>

Matrix




1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	1.8

Convolution  
Result

# Convolutional Layers

---

# Convolutional Layers



**Convolution layers - zoom in on specific bits of input**

**Extract structure and features in the input image**

**Successive layers aggregate inputs into higher level features**


**Pixels >> Lines >> Edges >> Object**




# Feature Maps



Image




Pixels




Feature  
Map

Feature maps are convolutional layers generated by applying a convolutional kernel to the input

# Feature Maps




Pixels




Convolutional  
Layer

Neurons




# Feature Maps



Number of neurons  
in receptive field =  
kernel size

# Feature Maps


Neuron i



Pixels

Convolutional  
Layer

# Kernel Size

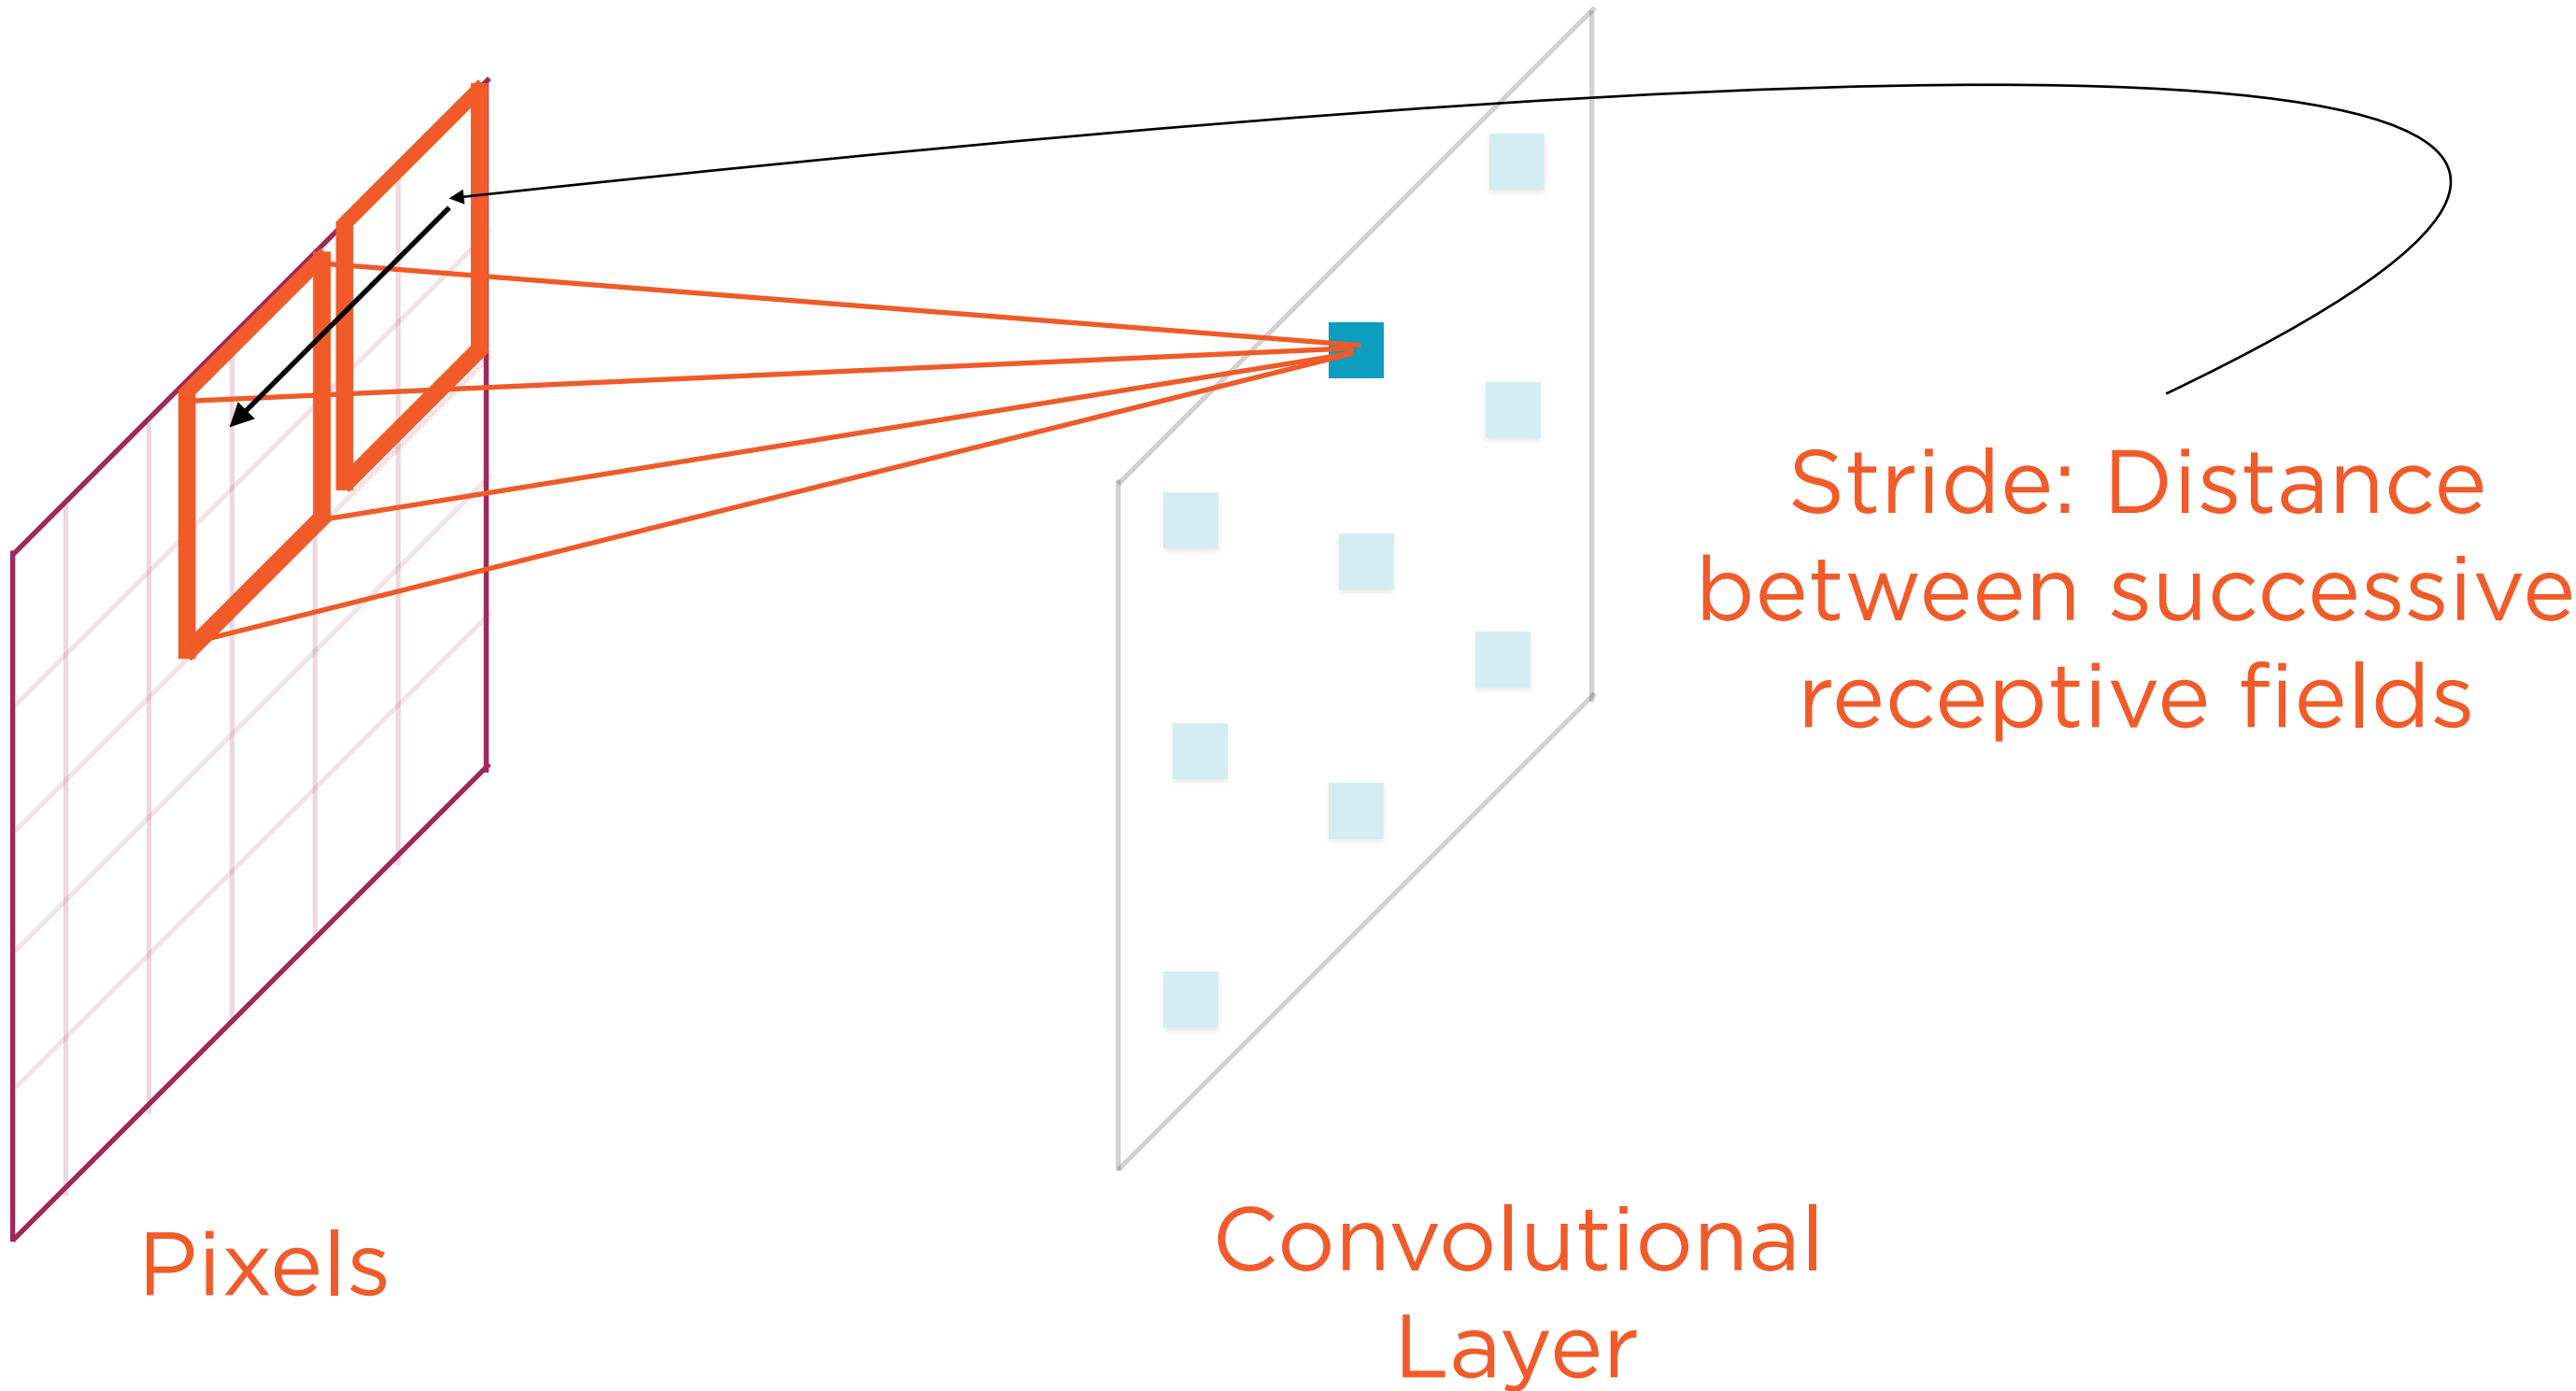


Convolutional kernel size usually expressed in terms of width and height of receptive area


Use **small convolutional kernels, more efficient**

Stacking two 3x3 kernels is preferable to one 9x9 kernel

# Feature Maps




# Feature Maps







# Feature Maps



# Feature Maps




# Feature Maps




# Feature Maps

Zero padding  
may be needed  
at the edges




# Feature Maps



Sparse, not  
Dense


# Feature Maps



**Notice also that neurons are not connected to all pixels**

**CNNs are *sparse* neural networks**

# Feature Maps




**All neurons in a feature map have the same weights and biases**

**Two big advantages over DNNs**

- Dramatically **fewer** parameters to train
- CNN can recognize feature patterns **independent** of location

# Feature Maps




The parameters of all neurons in a feature map are collectively called the filter

Why filter?


Because weights **highlight (filter)** specific patterns from the input pixels




# CNNs



Feature  
Map




Convolutional  
Layer



CNN

# Convolutional Layer




**Each convolutional layer consists of several feature maps of equal sizes**

**The different feature maps have different parameters**

# Pooling Layers

---


# Pooling



Matrix




Max,  
2x2 filter,  
stride = 2



Pooling Result


# Pooling



Matrix




Max,  
2x2 filter,  
stride = 2



Pooling Result


# Pooling



Matrix




Max,  
2x2 filter,  
stride = 2



Pooling Result


# Pooling



Matrix




Max,  
2x2 filter,  
stride = 2



Pooling Result


# Pooling



Matrix




Max,  
2x2 filter,  
stride = 2



Pooling Result



# Pooling Layers




**Neurons in a pooling layer have no weights or biases**

**A pooling neuron simply applies some aggregation function to all inputs**

**Max, sum, average**


# Pooling Layers



## Why use them?

- Greatly reduce memory usage during training
- Mitigate overfitting (via subsampling)
- Make NN recognize features independent of location (location invariance)

# Pooling Layers



**Pooling layers typically act on each channel independently**


**So, usually, output area < input area but**

**Output depth = Input depth**

# CNN Architectures


---

# Typical CNN Architecture




**Alternating groups of convolutional and pooling layers**

# Typical CNN Architecture




**Each group of convolutional layers usually followed by a ReLU layer**

# Typical CNN Architecture



**The output of each layer is also an image**


# Typical CNN Architecture



**However successive outputs are smaller and smaller (due to pooling layers)**




# Typical CNN Architecture



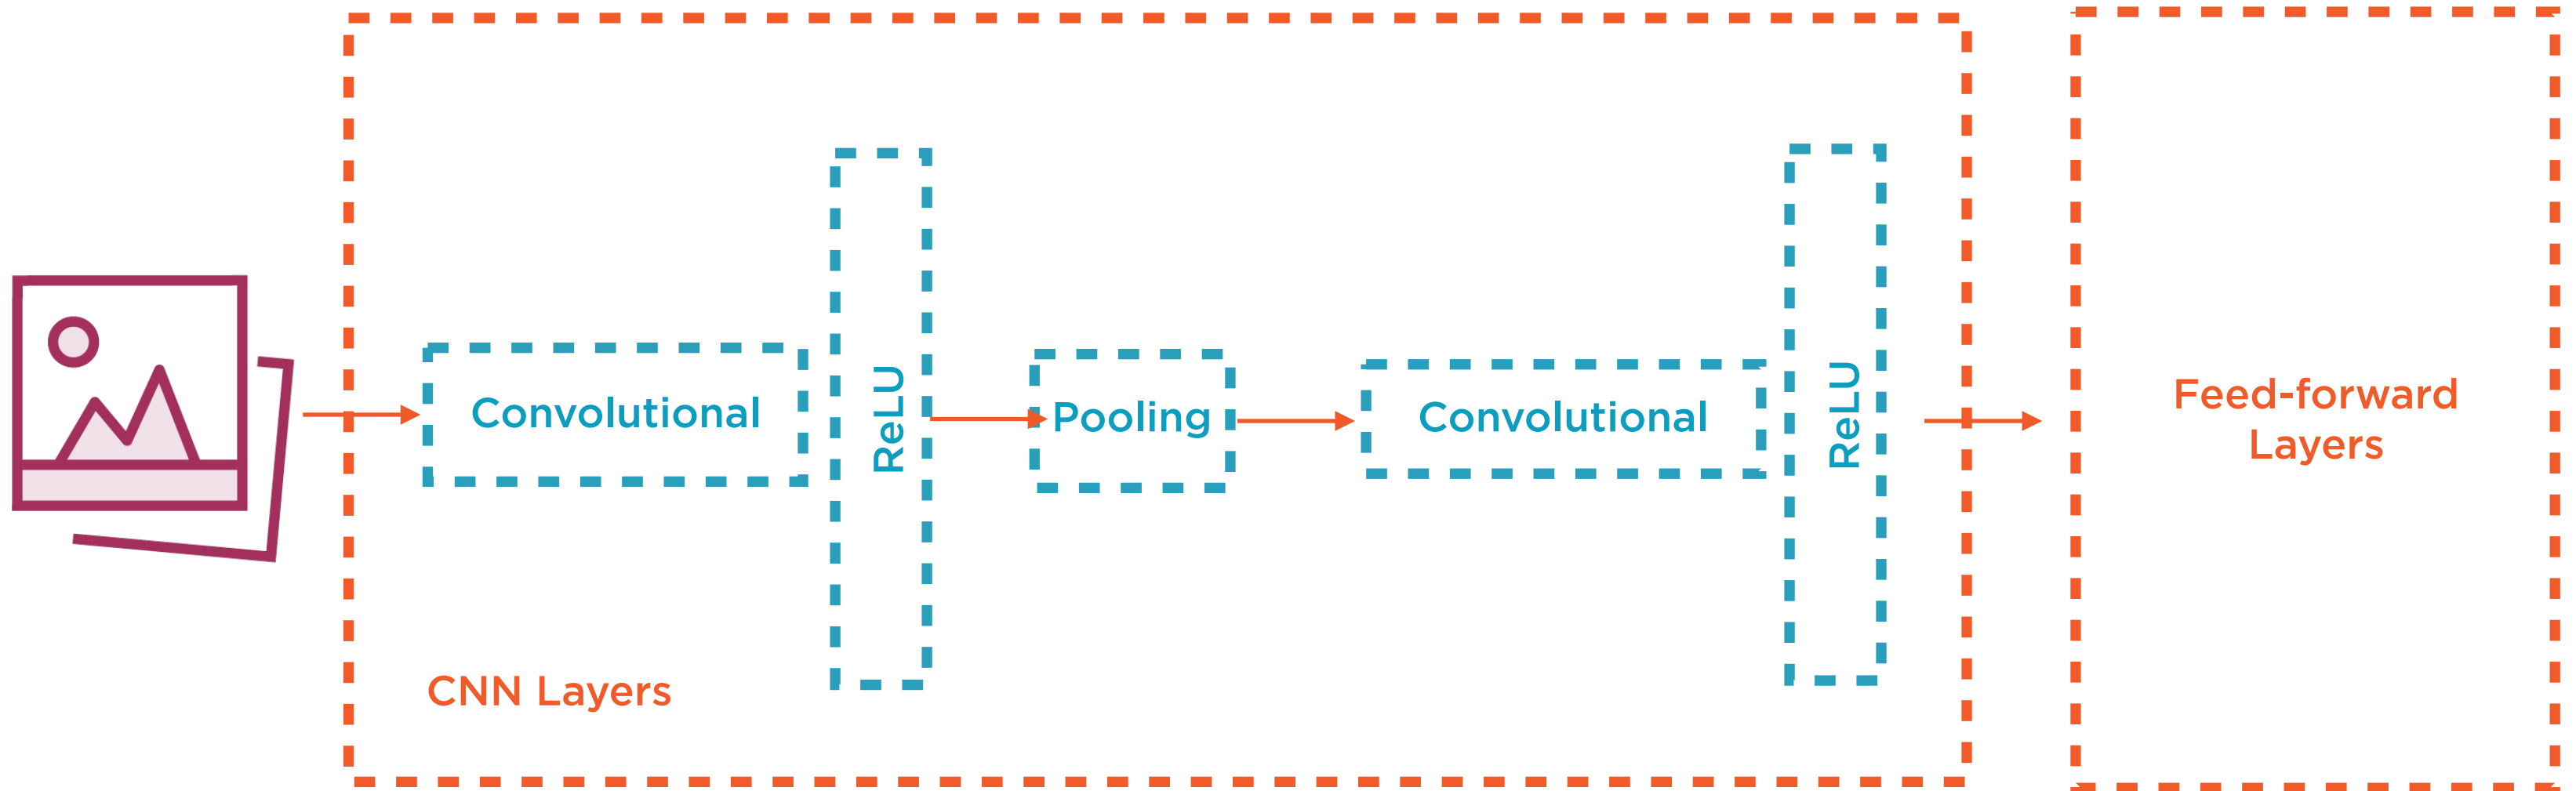
**As well as deeper and deeper (due to feature maps in the convolutional layers)**

# Typical CNN Architecture




**This entire set of layers is then fed into a  
regular, feed-forward NN**

# Typical CNN Architecture



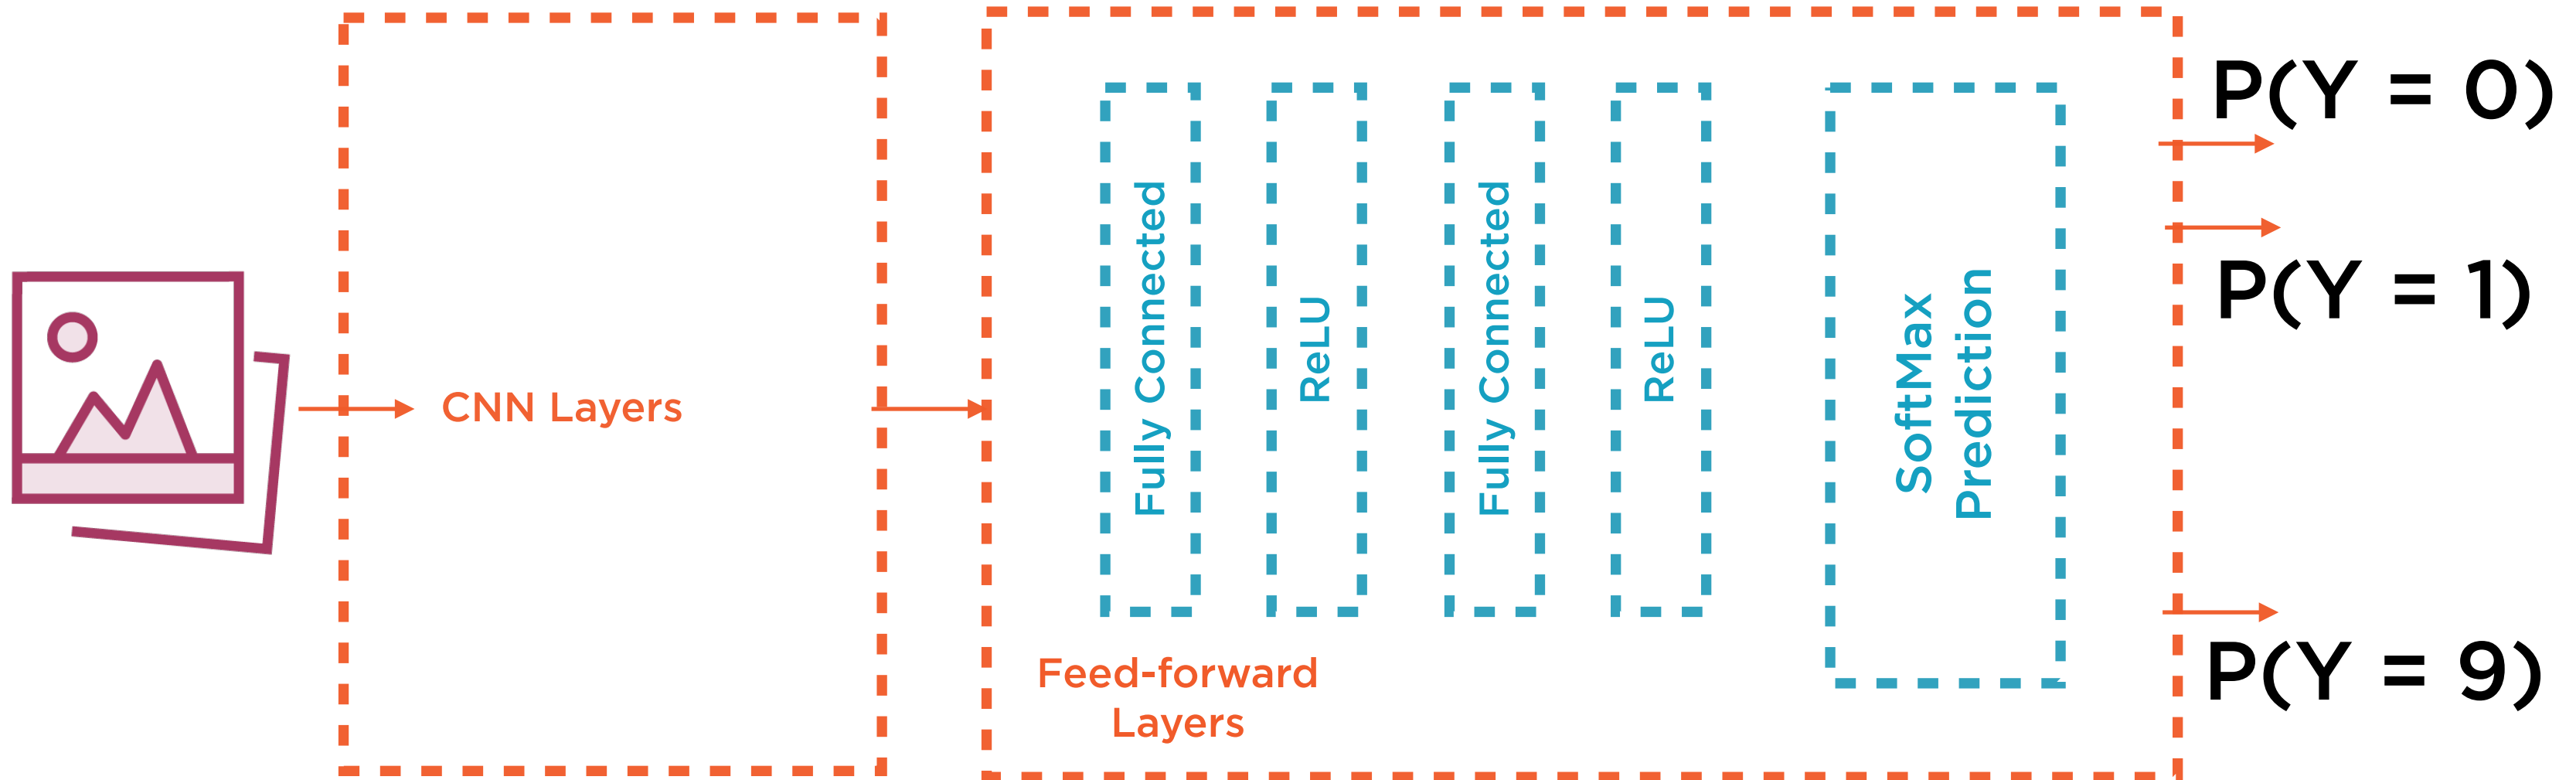
**This entire set of layers is then fed into a regular, feed-forward NN**

# Typical CNN Architecture




**This feed-forward has a few fully connected layers with ReLU activation**

# Typical CNN Architecture



**This is the output layer, emitting probabilities**

# Typical CNN Architectures



**Input is an image**

**Outputs are probabilities**

# Demo

**Apply convolution and pooling filters  
to images**

# Summary

**Intuition behind Convolutional Neural Networks (CNNs)**

**Convolution layers and feature maps**

**Pooling layers to subsample inputs**

**Typical CNN architecture**